

Agents that do the work. Not demos that *describe it*.

Matrion designs and builds production-grade agentic AI systems — grounded in your data, connected to your tools, evaluated before they ship, and operated by senior engineers embedded inside your team.

CORE FOCUS	DELIVERY MODEL	SHIPPED AS	MEASURED BY
Agentic systems	Embedded engineers	Production code	Evaluations

01 / WHAT WE BUILD

Agentic systems that move work forward.

We go beyond chat interfaces — agents that understand goals, plan work, call tools, collaborate with people and other agents, and complete auditable actions inside your existing operations.

01 Autonomous workflows Goal-driven agents that decompose complex work, coordinate steps and adapt — within clear boundaries and approvals.	02 Enterprise RAG & grounding Retrieval, citations, permissions and context controls that reduce hallucination and improve decisions.
03 Human-agent collaboration Approval gates, escalation paths and role-aware experiences that keep people in control where it matters.	04 Multi-agent orchestration Specialist agents that delegate, share context and coordinate across end-to-end business processes.
05 Evaluation & observability Traceable decisions, quality evaluations, cost and latency controls, and monitoring from day one.	06 Tools, actions & guardrails Safe connections to APIs and applications, with policy checks and failure handling that make actions dependable.
07 Generative Engine Optimization (GEO) Your customers increasingly ask an assistant instead of a search box. We optimise how your brand, products and documentation are retrieved, represented and cited inside AI-generated answers — and track the share of answers you actually appear in.	

40%+

of agentic AI projects will be cancelled by the end of 2027 — attributed to escalating costs, unclear business value and inadequate risk controls. **Model capability is not on that list.** These are delivery and governance failures, which is exactly where we work.

Source: Gartner, June 2025

Forward Deployed Engineers — inside your team, not across a contract.

Building dependable agents requires deep context about your workflows, systems and risk. We embed senior AI engineers directly into your company — in your repos, your data, your roadmap — until the agentic system is live and your team owns it.

01

Embedded from day one

Your stack, your tickets, your rituals — not an external vendor on a status call.

02

Ships production code, not decks

Agents and evaluations that land in your repository and run against real traffic.

03

Weeks to first value

A narrow, real use case goes live early while the wider roadmap is still forming.

04

Built to hand over

Knowledge transfer is the deliverable. Scale us up or down as you need.

How an engagement actually runs.

WEEK 0

Scope and baseline

One workflow with a real owner. The evaluation set is built before the agent, so “better” is defined first.

WEEKS 1-3

Thin slice to production

The narrowest useful version against real traffic, behind approval gates, with tracing and rollback.

WEEKS 4-10

Harden and widen

Coverage expands as evaluations hold. Failure handling and permissions grow with the surface area.

ONGOING

Operate and hand over

Your engineers take ownership with the harness, runbooks and evaluation suite intact.

Want your website optimised for GEO?

Your customers increasingly ask an assistant instead of a search box. We audit how your brand, products and documentation are retrieved, represented and cited inside AI-generated answers — then fix the structure, sourcing and machine-readability so you show up. Ask us for a GEO audit.

Built for the deadlines that are actually coming.

IN	DPDP Act – 14 May 2027	India's substantive obligations — notice, consent, safeguards, breach reporting, data principal rights — with consent-manager provisions from November 2026.
IN	Synthetic content labelling	The IT Amendment Rules 2026 have applied since 20 February 2026, requiring synthetically generated information to carry visible or audible labels.
EU	EU AI Act – 2 Dec 2027	The Digital Omnibus moved standalone high-risk obligations from August 2026 to December 2027; transparency and marking duties apply from December 2026.
STD	ISO/IEC 42001 readiness	Control evidence — evaluation records, model and data lineage, human oversight — produced as a by-product of delivery, not a retrofit before an audit.

Models	Frontier tier – Claude / GPT / Gemini · Open-weight – Llama, Qwen, DeepSeek, Mistral – for self-hosting and residency · Routed per task
Agent runtimes	Claude Agent SDK · OpenAI Agents SDK · Microsoft Agent Framework · Google ADK · Purpose-built harnesses
Interoperability	MCP – agent to tools and data · A2A – agent to agent · Both under the Linux Foundation
Context engineering	Retrieval (RAG) · Compaction · Tool-result clearing · Long-horizon memory · Permission-aware context
Evaluation	LangSmith · Braintrust · Langfuse · Arize Phoenix · W&B Weave · OpenTelemetry
Serving & platform	vLLM · AWS Bedrock · Google Vertex AI · Azure AI Foundry · Prompt caching and batch

Let's put agents to work
on your hardest workflow.

EMAIL

info@matrion.in

WEB

matrion.in

PHONE

+91 98455 76009

FOCUS

Agentic systems